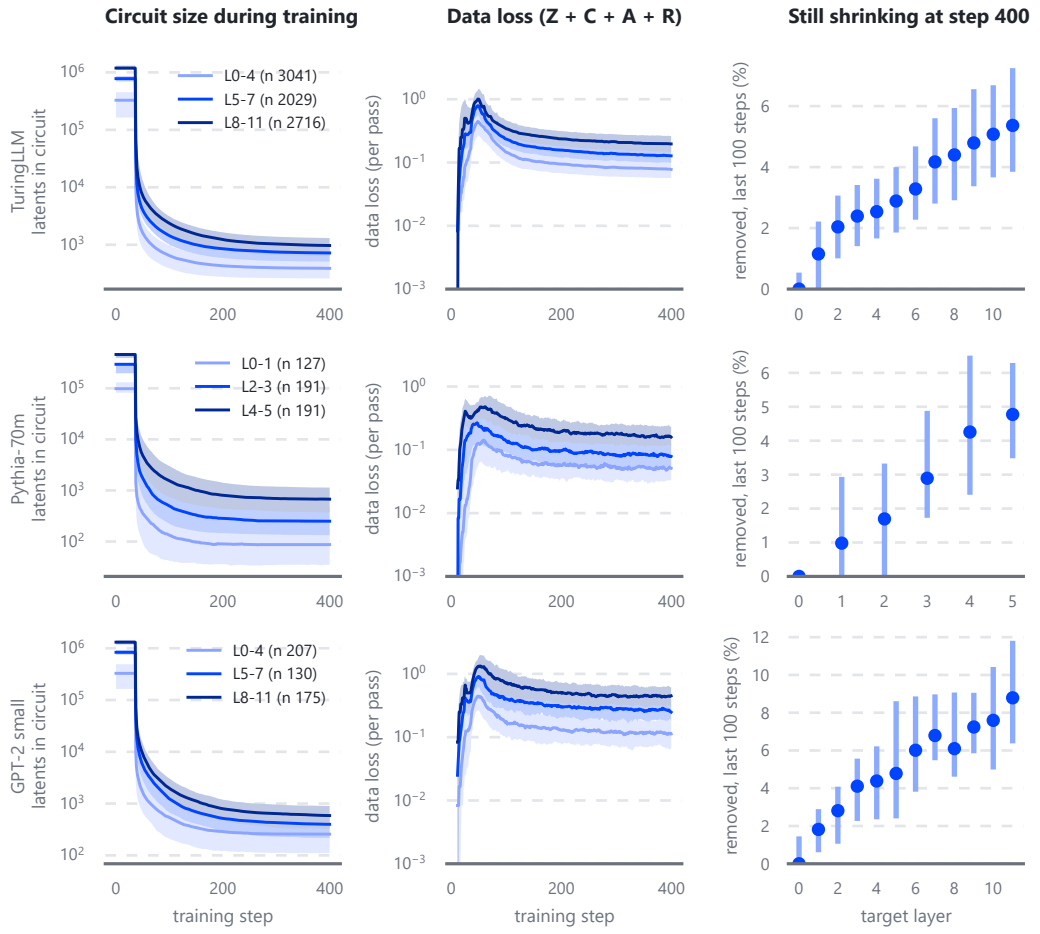

Circuit size during training
Data loss (Z + C + A + R)
Still shrinking at step 400
TuringLLM
latents in circuit
L0-4 (n 3041)
L5-7 (n 2029)
L8-11 (n 2716)
Pythia-70m
latents in circuit
L0-1 (n 127)
L2-3 (n 191)
L4-5 (n 191)
GPT-2 small
latents in circuit
L0-4 (n 207)
L5-7 (n 130)
L8-11 (n 175)
data loss (per pass)
removed, last 100 steps (%)
training step
target layer