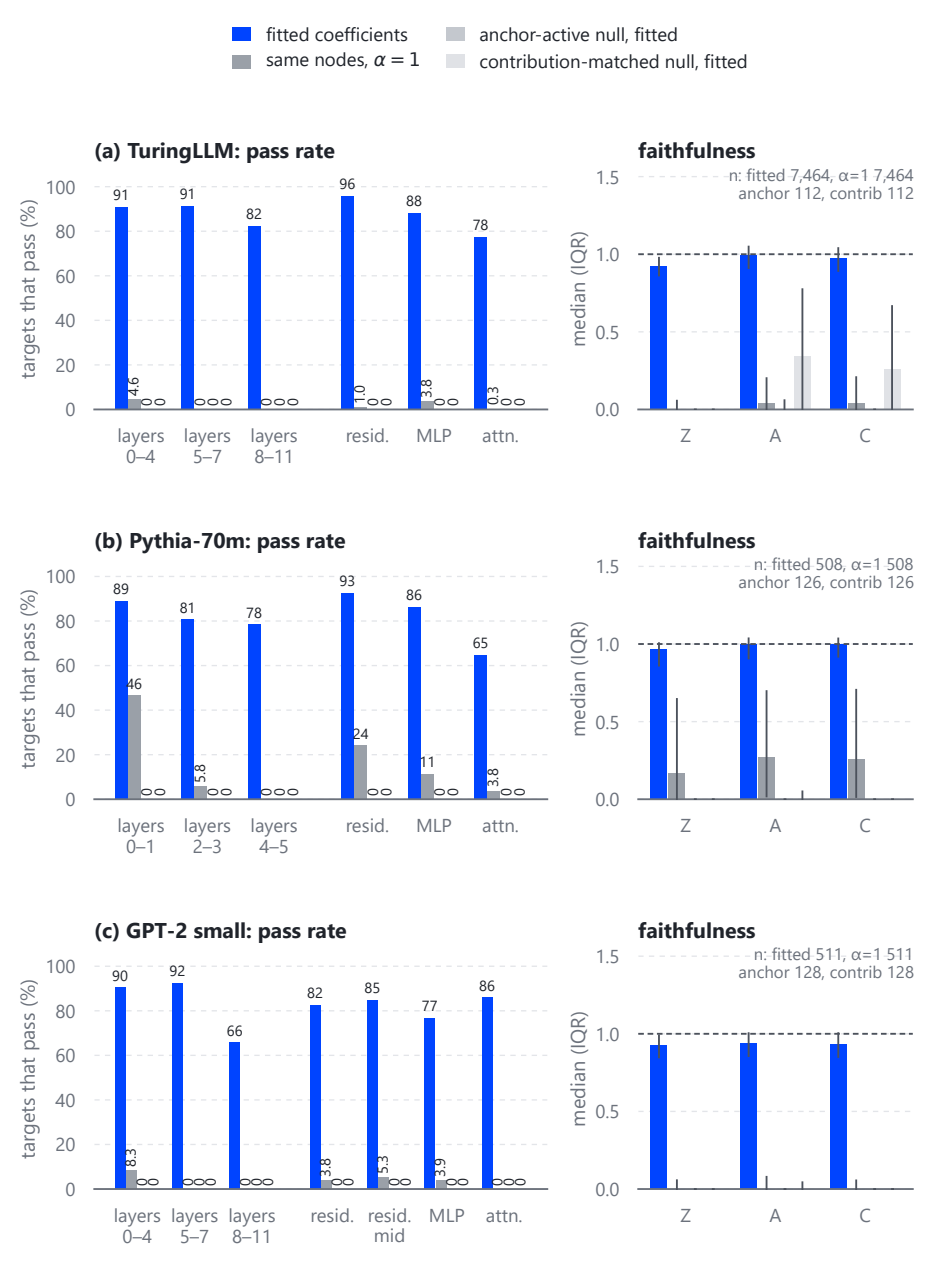
fitted coefficients
same nodes, α = 1
anchor-active null, fitted
contribution-matched null, fitted
(a) TuringLLM: pass rate
targets that pass (%)
91
4.6
0
0
91
0
0
0
82
0
0
0
96
1.0
0
0
88
3.8
0
0
78
0.3
0
0
layers 0–4
layers 5–7
layers 8–11
resid.
MLP
attn.
faithfulness
n: fitted 7,464, α=1 7,464
anchor 112, contrib 112
median (IQR)
Z
A
C
(b) Pythia-70m: pass rate
89
46
0
0
81
5.8
0
0
78
0
0
0
93
24
0
0
86
11
0
0
65
3.8
0
0
layers 0–1
layers 2–3
layers 4–5
n: fitted 508, α=1 508
anchor 126, contrib 126
(c) GPT-2 small: pass rate
90
8.3
0
0
92
0
0
0
66
0
0
0
82
3.8
0
0
85
5.3
0
0
77
3.9
0
0
86
0
0
0
layers 0–4
layers 5–7
layers 8–11
resid.
resid. mid
n: fitted 511, α=1 511
anchor 128, contrib 128